

Introduction to Machine Learning for Bioinformatics

Sonika Tyagi, Navya Tyagi and Tyrone Chen
RMIT University Australia

Main contributor to Big data growth is genomics






Fig by Tyrone Chen @tyagilab using data from Stephens et al, 2015

REGULATORY OMICS SIGNATURES DRIVE FUNCTIONAL OMICS SIGNATURES

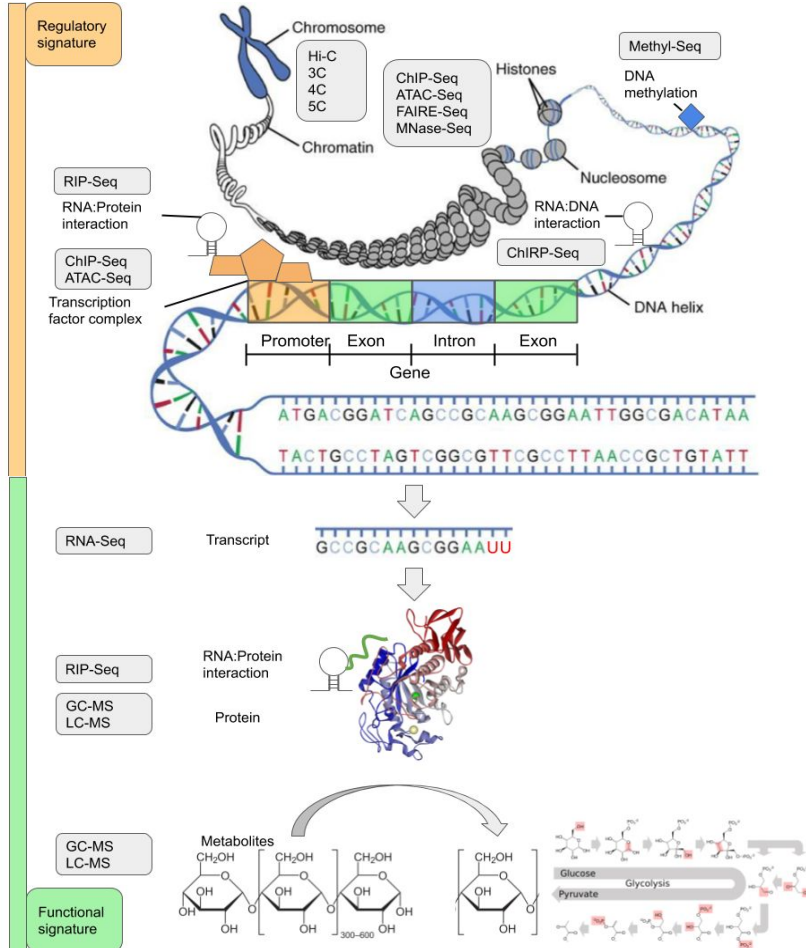
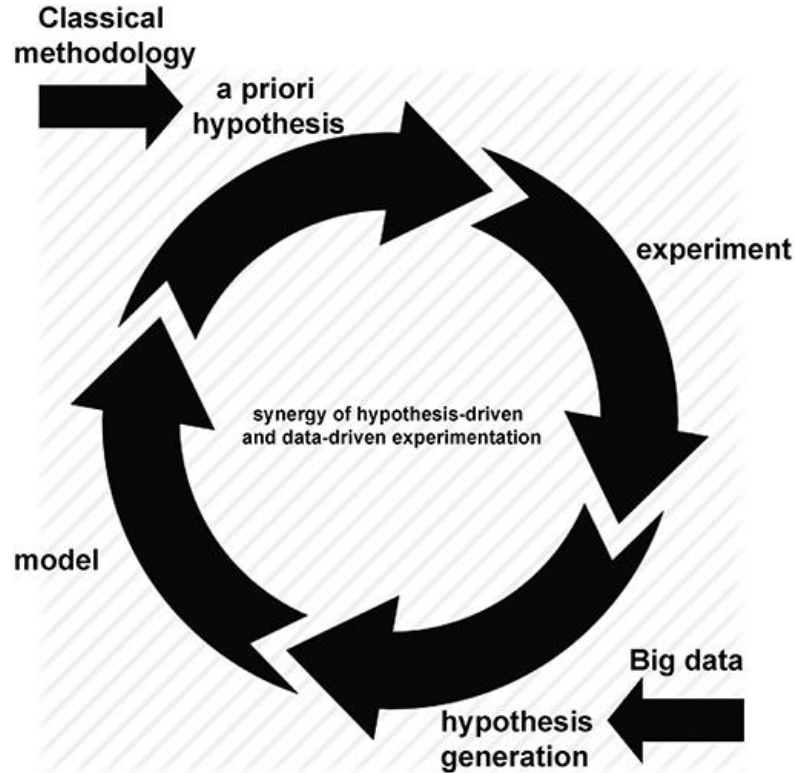


Image source: TyagiLab 2020

Genomics as a data-driven discipline



<https://doi.org/10.3389/fmed.2019.00034>


What is machine learning (ML) and Artificial intelligence (AI)?

AI = making intelligent machines

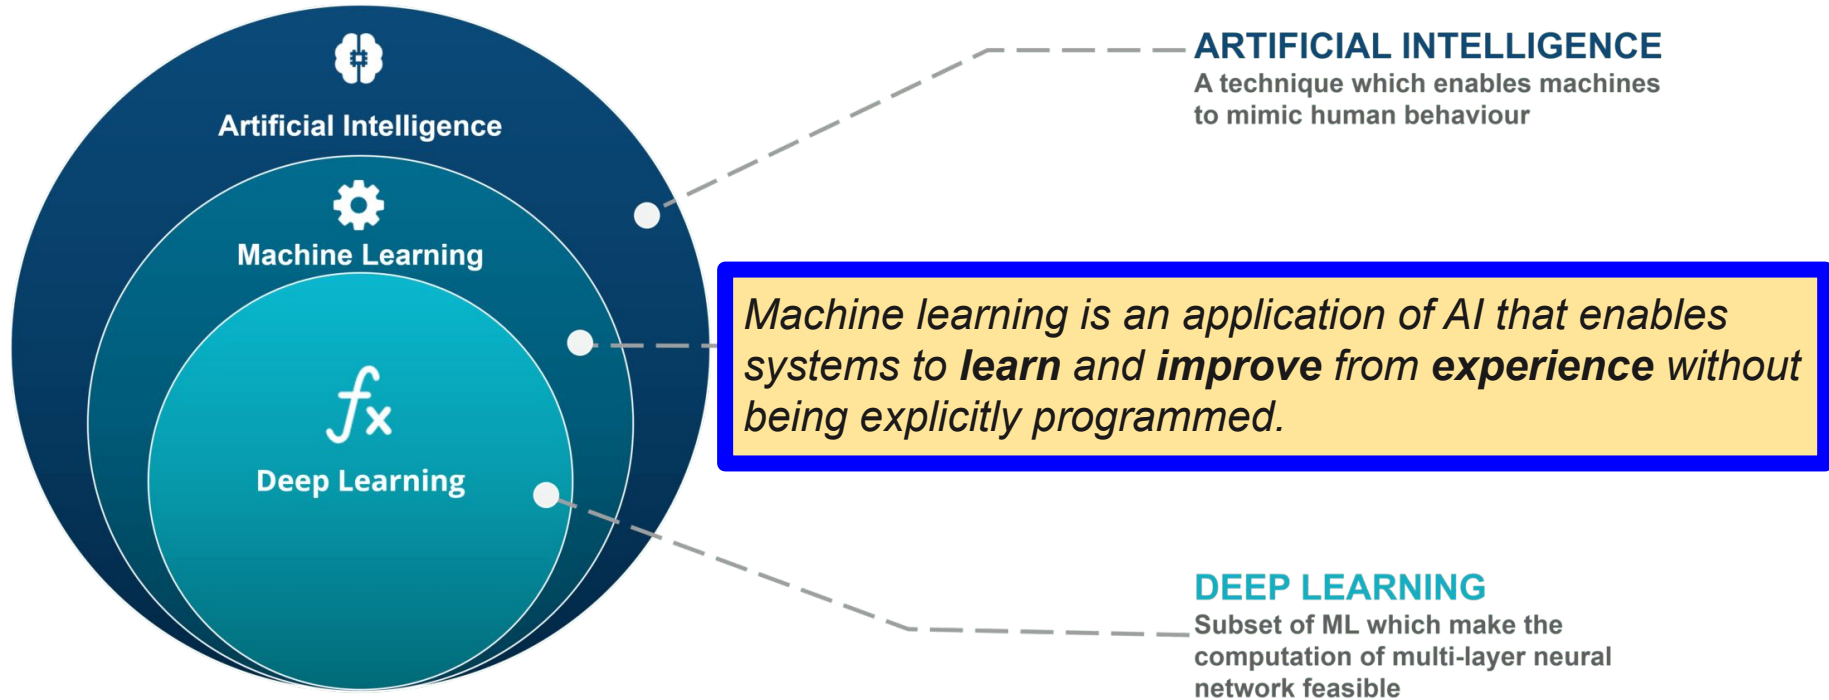
“Machines are intelligent to the extent that their actions can be expected to achieve their objectives”

-Prof Stuart Russell

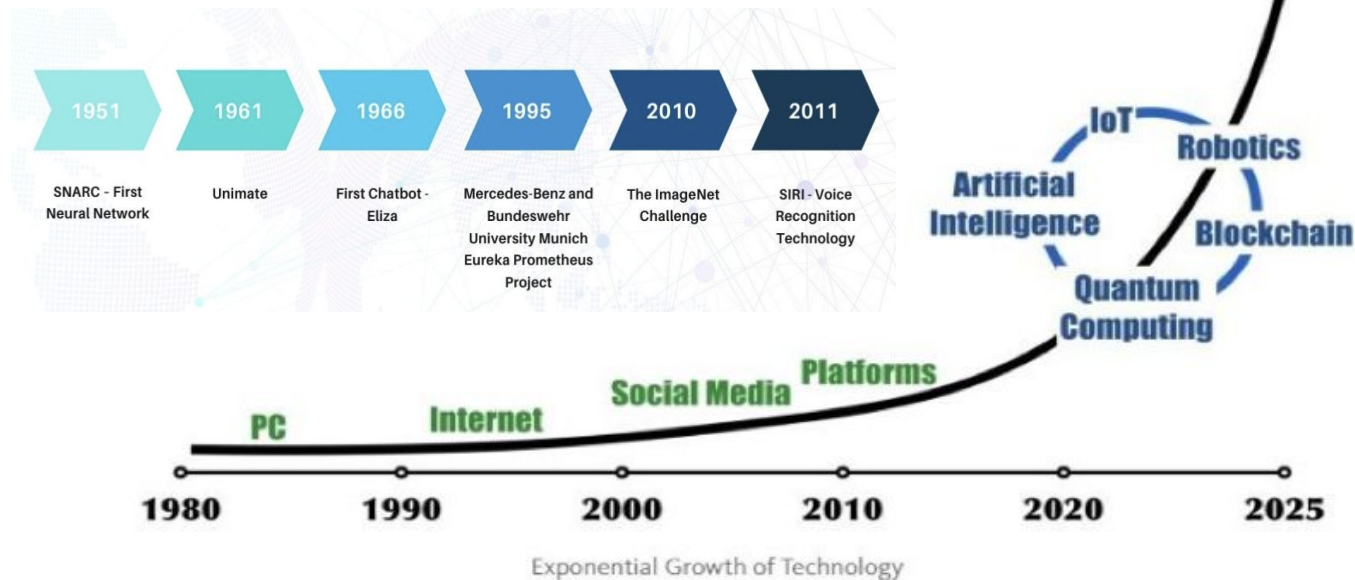
Machine learning (ML) ?



Machine learning (ML) ?



Evolution of AI



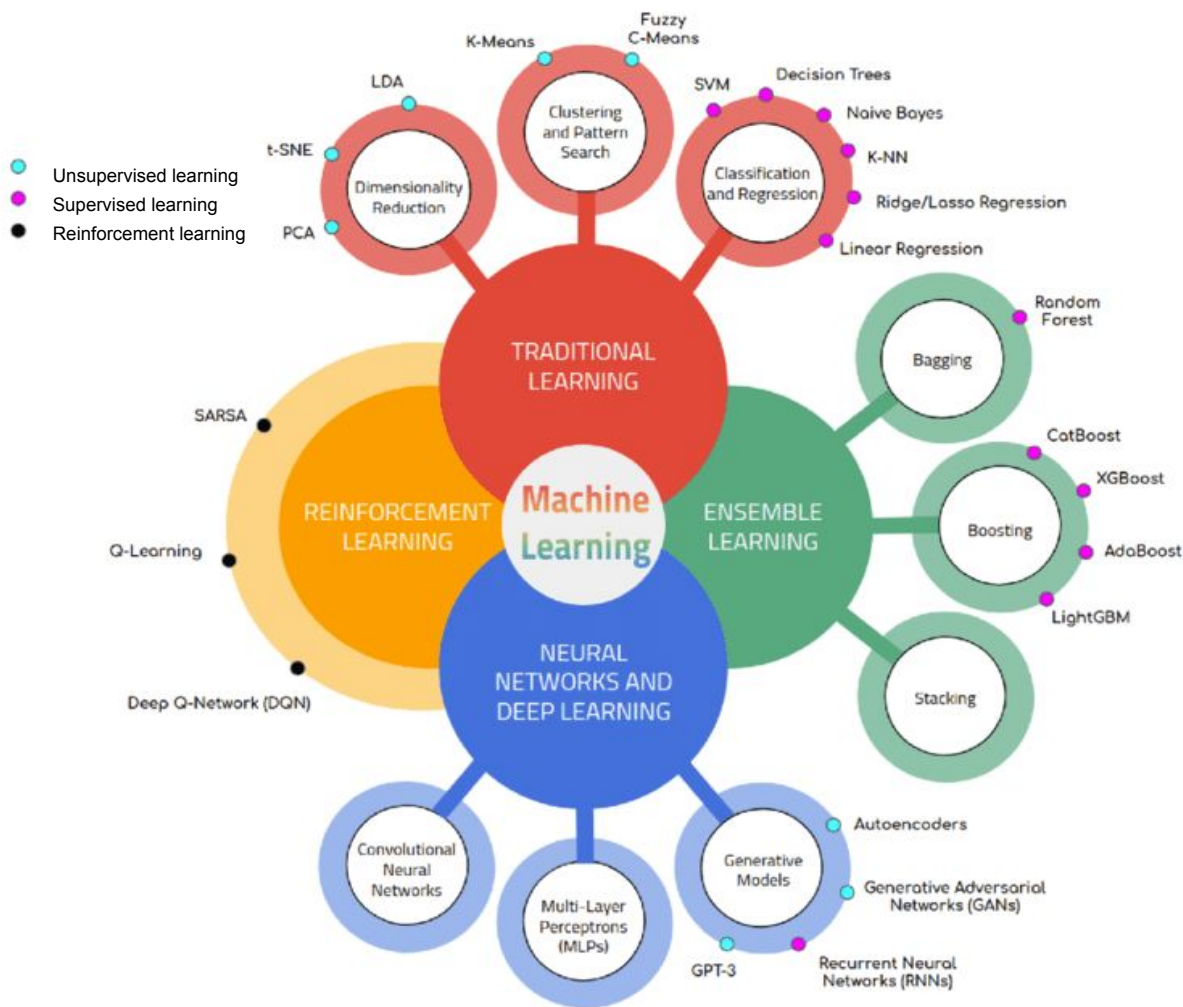



Image source:

DOI: [10.1007/s12039-021-01995-2](https://doi.org/10.1007/s12039-021-01995-2)


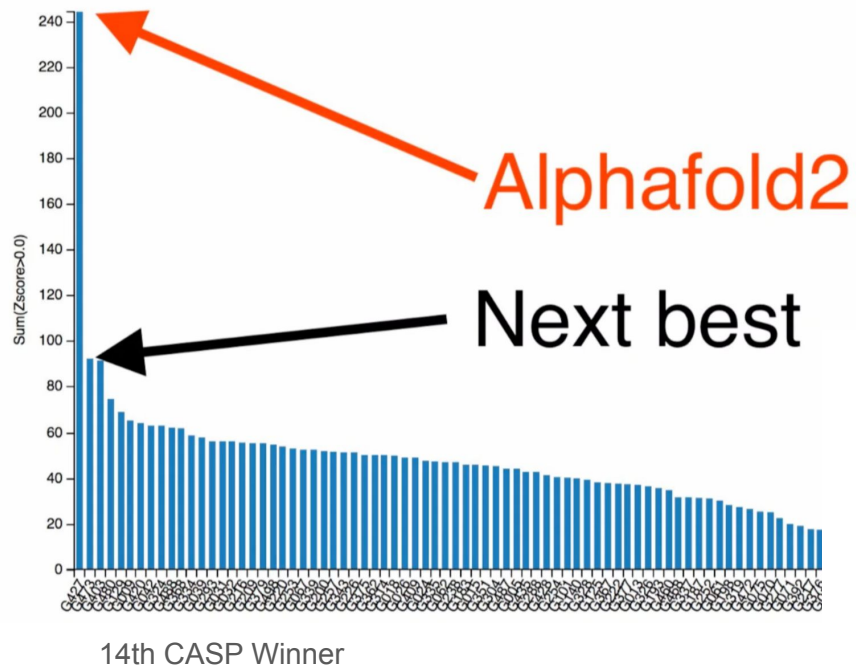
Deep Learning Applications for Genomics

1. Pattern recognition
2. Predicting biomolecule structures
3. Classification or predictive modeling
4. Image analysis




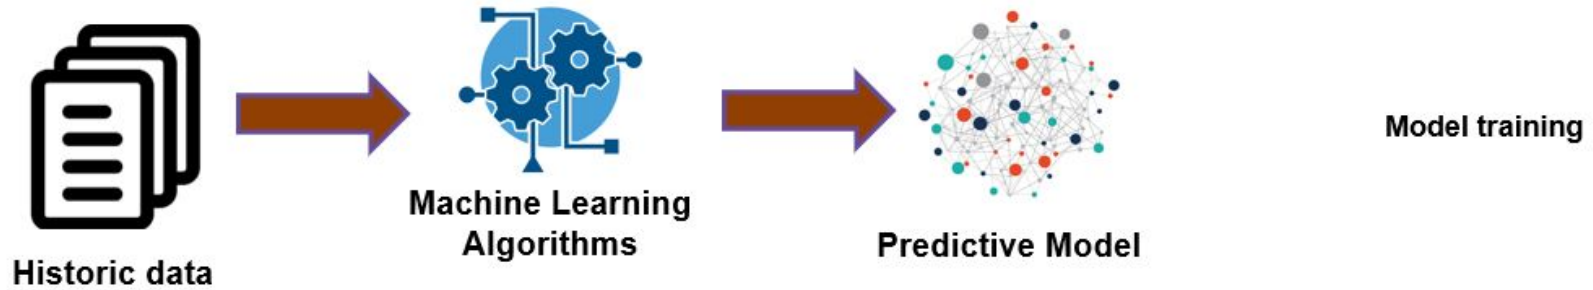
+
Transformers and
autoencoders

Eraslan *et al* 2019




A typical ML workflow: Supervised Learning

How ML modeling is done?



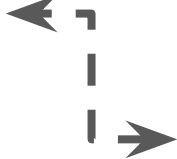
Labelled data for supervised learning

	seq_id	sequence	labels
0	NULL_F_sacCer3_ct_tbnabiRefSeq_7481_NM_0011845...	GTCTACCTACCTTATTAAGATCTGGGGATTAGAGCGGAGCAGCACC...	0
1	RC__NULL_F_sacCer3_ct_tbnabiRefSeq_7481_NM_001...	TACAGTCCAAGCGGACTCATGTCTGATTCATATCACAAAGGCTTGGT...	0
2	NULL_F_sacCer3_ct_tbnabiRefSeq_7481_NM_0011800...	TGTCCAATTGTAATCAATTCATGGGTCAAGAATAACGGTTCATTGT...	0
3	RC__NULL_F_sacCer3_ct_tbnabiRefSeq_7481_NM_001...	TACGCATACGCCTCAGTATAGCAATTAGGCAGCTTTGTTGCACAGT...	0
4	NULL_F_sacCer3_ct_tbnabiRefSeq_7481_NM_0011782...	ATCCTAAATACTCCGTTGTAAAGCATGTAAATAATATGCACAAAGC...	0



Training and test data: Dividing into two or three parts:

Training data: train the model




Test data: test and iterate until we have the best model

The diagram consists of a vertical dashed line. At the top of the line, a solid arrow points left towards the 'Training data' text. At the bottom of the line, a solid arrow points right towards the 'Test data' text. This visualizes an iterative process where the model is trained and then tested, with the results of the test used to refine the training.

Validation data: validate on a held out unseen data

What happens during training:




Training a model simply means **learning** i.e. determining good values (for all the **weights** and the **bias**) from labeled examples.

Choosing parameters that minimize the loss

Can we have a direction to go in a parameter space?

Compute the gradient derivative of loss function





Bias and Variance

We want to keep error as low as possible

Two major sources of error are bias and variance.

Assumption of supervised learning:

- There is a relationship between the feature(s) and target
- We estimate this relationship using a model (unknown)
- We train different training set to build a model and measure (repeat)
 - The difference between the outcomes of these models describing the relationship between features and target is called “variance”
 - Assumption about the relationship between feature(s) and target is simplified as “bias”




Decision trees

- The decision tree classifiers organizes a series of test questions and conditions in a tree structure.
- The root and internal nodes contain attribute test conditions to separate data entries that have different characteristics.
- All the terminal node is assigned a class label Yes or No.


A

Strain ID	Input features		Response
	Protein 1 (units/cell)	Protein 2 (units/cell)	Production (Low/High)
s1	2	8	Low
s2	3	6	Low
s3	9	6	Low
s4	5	10	High
s5	10	10	Low
s6	5	1	Low
s7	6	2	Low
s8	9	1	Low
s9	6	7	High
s10	1	1	Low
s11	5	4	Low
s12	2	4	Low
s13	6	12	High
s14	9	3	Low
s15	4	7	?

C



B



Decision trees pros and cons

- ❖ Scale invariant
- ❖ Robust to irrelevant features
- ❖ Interpretability
- ❖ Prone to overfitting
- ❖ Don't generalize well
- ❖ Pruning can help alleviate but can not diminish the effects

Random Forest

<i>id</i>	x_0	x_1	x_2	x_3	x_4	y
0	4.3	4.9	4.1	4.7	5.5	0
1	3.9	6.1	5.9	5.5	5.9	0
2	2.7	4.8	4.1	5.0	5.6	0
3	6.6	4.4	4.5	3.9	5.9	1
4	6.5	2.9	4.7	4.6	6.1	1
5	2.7	6.7	4.2	5.3	4.8	1


<i>id</i>
2
0
2
4
5
5

<i>id</i>
2
1
3
1
4
4

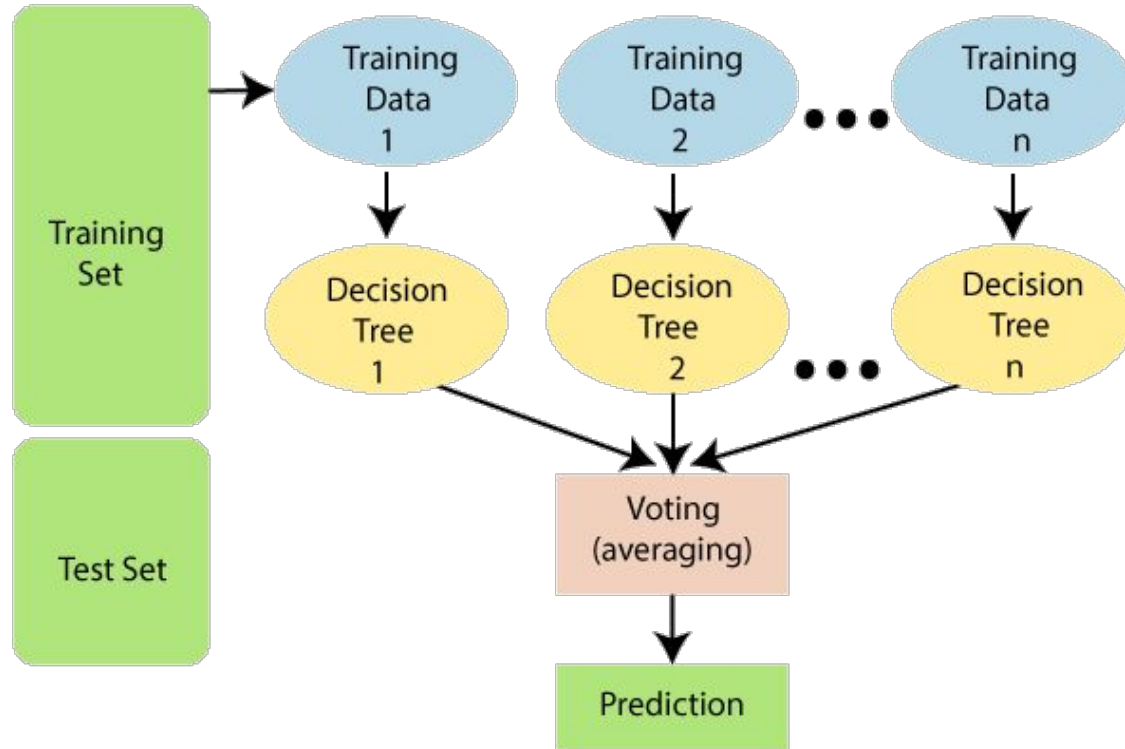
<i>id</i>
4
1
3
0
0
2

<i>id</i>
3
3
2
5
1
2

Bagging =
Bootstrapping +
Aggregating



Workflow:



The dataset is divided into subsets and given to each decision tree. During the training phase, each decision tree produces a prediction result, and when a new data point occurs, then based on the majority of results, the Random Forest classifier predicts the final decision.

Decision trees pros and cons

- ❖ Scale invariant
- ❖ Robust to irrelevant features
- ❖ Interpretability
- ❖ Prone to overfitting
- ❖ Don't generalize well
- ❖ Pruning can help alleviate but can not diminish the effects

RF: Thousands of trees (each tree may overfit but their combined decision is considered by majority votes)

-Allowing variability in both dimensions allows better generalization.

Confusion matrix




Source: <https://medium.com/@neeraj.kumar.iitg/statistical-performance-measures-12bad66694b7>

Confusion matrix:

		Predicted Class		
		Positive	Negative	
Actual Class	Positive	True Positive (TP)	False Negative (FN) Type II Error	Sensitivity $\frac{TP}{(TP + FN)}$
	Negative	False Positive (FP) Type I Error	True Negative (TN)	Specificity $\frac{TN}{(TN + FP)}$
		Precision $\frac{TP}{(TP + FP)}$	Negative Predictive Value $\frac{TN}{(TN + FN)}$	Accuracy $\frac{TP + TN}{(TP + TN + FP + FN)}$

ML model performance assessment



Confusion matrix

Accuracy

F1 Score

Precision

Recall

FPP = 1-specificity

TPP=sensitivity

Preprocessing of Genomic Data

Slides by Naima Vahab

NLP preprocessing of sequence data:

1. Tokenization: Finding K-mers
2. Numerical encoding:
 - a. Ordinal Encoding
 - b. One-Hot Encoding
3. Transformations
 - a. Frequency tables
 - i. N-grams
 - ii. Term Frequency - Inverse Document Frequency
 - b. Embedding Vector
 - c. Positional encoding

Terminology used in this document:

Token: smallest unit of processing text data. Also known as k-mer in genomics.

Corpus: reference library or database of all tokens

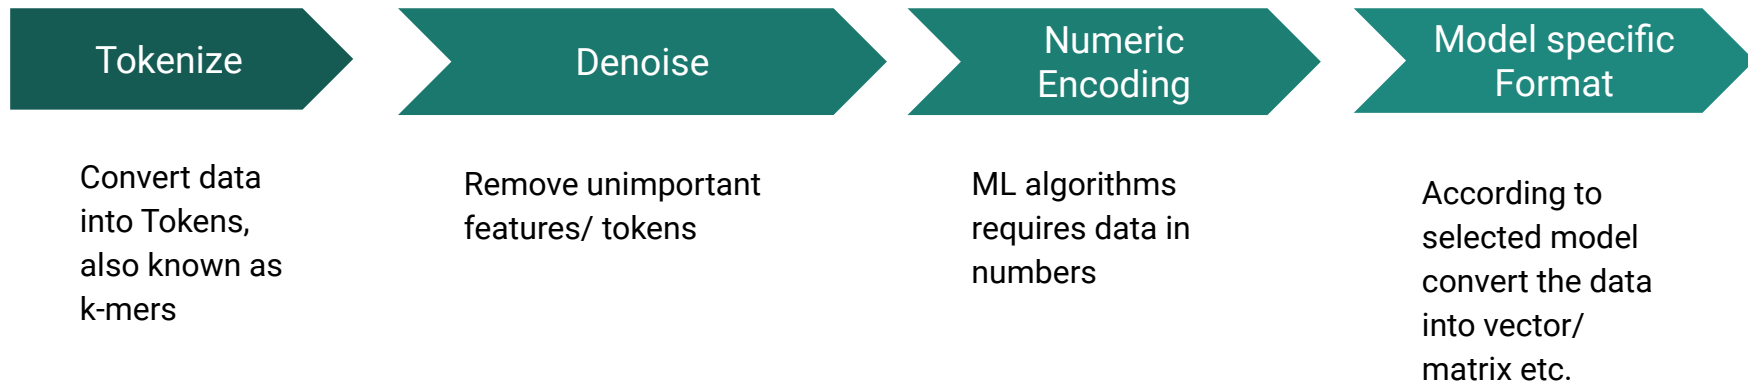
Vector: a set of numbers

Matrix: a two dimensional table of numbers

Stop_words: token with less information content e.g. the, a, an etc in English or known repeat patterns in a genomic sequence

Encoding: converted into a coded form

Workflow : Before applying ML models on genomic sequence, following steps are used to preprocess the data



Tokenization : Split the genome sequence in arbitrary length

Example:

INPUT : TAATG...

KMER_1: TAA--

KMER_2: -AAT-

KMER_3: --ATG

```
def get_kmers(sequence: str, length: int):  
    """  
    Take a dna sequence as input and split the sequence into k-mers / tokens  
    """  
    return [sequence[x:x+length].upper() for x in range(len(sequence) - length + 1)]  
  
# in this example, we will split on length 3 k-mers  
length = 3  
kmers = get_kmers(input_sequence, length)  
print(kmers)
```

Denoise the data : Removing less important features or stop words

```
def filter_kmers(tokens: str, stopwords: list):  
    """  
    Take an input dna sequence and list of stopwords to remove from the data.  
    """  
    return [x for x in tokens if x not in stopwords]  
  
# in this example, let us pretend this list of k-mers have low information content  
stopwords = ["TAA", "AAT"]  
filtered_kmers = filter_kmers(kmers, stopwords)  
print(filtered_kmers)
```


Ordinal encoding : Assigns an integer to each category value

```
from sklearn.preprocessing import OrdinalEncoder

def encode_ordinal(input_sequence: str, length: int):
    """
    Take a list of k-mers and perform ordinal encoding
    """
    tokens = [[x] for x in get_kmers(input_sequence, length)]
    encoder = OrdinalEncoder()
    return encoder.fit_transform(tokens)

ordinal = encode_ordinal(input_sequence, length)
print(ordinal)
```

```
[[5.]
 [0.]
 [1.]
 [7.]
 [4.]
 [3.]
 [2.]
 [6.]]
```

```
input_sequence = "AATCGAAAAAAAAA"  
output = ordinalen(input_sequence,length)  
print(output)
```

```
[[1.]  
 [2.]  
 [5.]  
 [3.]  
 [4.]  
 [0.]  
 [0.]  
 [0.]  
 [0.]  
 [0.]  
 [0.]]
```

One-Hot Encoding : Each label is represented by binary format

```
from numpy import asarray
from sklearn.preprocessing import OneHotEncoder

def onehoten(sequence, size):
    r"""Array of categories(here string) are sorted and returns binary variables for
    doc=sentence_in_list(sequence, size)
    data = asarray(doc)
    encoder = OneHotEncoder(sparse=False)
    onehot = encoder.fit_transform(data)
    return onehot
onehoten(input_sequence, length)
```

```
[[ 0 0 0 0 0 1 0 0]
 [ 1 0 0 0 0 0 0 0]
 [ 0 1 0 0 0 0 0 0]
 [ 0 0 0 0 0 0 0 1]
 [ 0 0 0 0 1 0 0 0]
 [ 0 0 0 1 0 0 0 0]
 [ 0 0 1 0 0 0 0 0]
 [ 0 0 0 0 0 0 1 0]]
```


TF-IDF (Term Frequency-Inverse Document Frequency) : Multiplying TF and IDF implies a weight to term which gives how important is a word in the document.

TF = (Number of repetitions of word in a document) / (# of words in a document)

IDF = $\text{Log}[(\text{Number of documents}) / (\text{Number of documents containing the word})]$

Words	IDF Value		Words/ Documents	Document 1	Document 2	Document 3
going	0		going	0.16	0.16	0.12
to	0.41		to	0.16	0	0.12
today	0.41		today	0.16	0.16	0
i	0.41		i	0	0.16	0.12
am	0.41		am	0	0.16	0.12
It	1.09		It	0.16	0	0
is	1.09		is	0.16	0	0
rain	1.09		rain	0.16	0	0

IDF Value and TF value of 3 documents.

```
from sklearn.feature_extraction.text import TfidfVectorizer
```

```
def tfidf(input_sequence: str, length: int):
```

```
    """
```

```
    Take a dna sequence and k-mer size as input,  
    output a matrix of TF-IDF features.
```

```
    """
```

```
    data = [" ".join(get_kmers(input_sequence, length))]
```

```
    tfidf = TfidfVectorizer()
```

```
    tfidf = tfidf.fit_transform(data)
```

```
    return tfidf.toarray()
```

```
tfidf_matrix = tfidf(input_sequence, length)
```

```
print(tfidf_matrix)
```

```
input_sequence = "AATCGAAAAAAAA"  
output = tfidf(input_sequence,length)  
print(output)
```


tf-idf values in matrix form:

```
[[0.93704257 0.15617376 0.15617376 0.15617376 0.15617376 0.15617376]]
```

Embedding Vector :

Taking a single sequence and projecting into the embedding returns a unique vector for each sequence.

Can use libraries like word2vec, dna2vec, GloVe etc.




```
import gensim
from gensim.models import Word2Vec

def create_word2vec(input_sequence: str, length: int):
    """
    Input a list of tokens to generate an embedding of word vectors
    """
    data = get_kmers(input_sequence, length)
    return Word2Vec([data], min_count=1)

def map_token(input_sequence: str, model: gensim.models.base_any2vec):
    """
    Input a token and embedding and map the token onto the embedding
    """
    return model[input_sequence]

embedding = create_word2vec(input_sequence, length)
vector = map_token('TAA', embedding)
print(vector)
```

```
input_sequence = "AATCGATAA"  
word2vec(input_sequence,length)
```

```
WARNING:gensim.models.base_any2vec:under 10 jobs per worker:  
[ 3.6882360e-03 -6.3672276e-05  3.1452794e-03  6.4240553e-04  
 -4.6913270e-03  9.3815307e-04  1.8332954e-04 -2.9308046e-03  
  8.7845704e-04 -1.5818959e-03  2.3677319e-03 -4.0797507e-03  
  4.4199773e-03 -3.2845191e-03 -2.4677652e-03  4.2411815e-03  
 -4.3340065e-03  4.4469675e-03 -4.8701679e-03  3.4830945e-03  
 -2.4540696e-04  4.7534686e-03  2.8795649e-03 -2.0364951e-03  
  1.4299533e-03  2.1512657e-03  1.1004285e-03 -3.7372194e-03  
 -1.5636545e-03  5.9850125e-05 -2.3654576e-03  3.5608963e-03  
 -4.8642256e-03 -3.2035201e-03  2.5428815e-03  2.1336516e-03  
 -3.2975324e-03  1.2050702e-03  2.5439151e-03 -4.2483918e-03]
```

Models in Summary

Input Sequence : AATAAGTGC

Ordinal Encoding : [[1],[2],[0],[3]]

One-Hot Encoding : [[0,1,0,1],[0,0,1,0],[1,0,0,0],[0,0,0,1]]


TF-IDF : [[0.93704257 0.15617376 0.15617376 0.15617376 0.15617376 0.15617376]]

**Embedding Vector : [3.6882360e-03 -6.3672276e-05 3.1452794e-03 6.4240553e-04
-4.6913270e-03 9.3815307e-04 1.8332954e-04 -2.9308046e-03
8.7845704e-04 -1.5818959e-03 2.3677319e-03 -4.0797507e-03
4.4199773e-03 -3.2845191e-03 -2.4677652e-03 4.2411815e-03]**

Introducing **genomeNLP** by Tyrone Chen et al 2022-23

MODERN DEEP LEARNING TOOLKIT FOR BIOLOGICAL DATA

1 Problem



High barrier for biologists

Existing high-level machine learning interfaces are tailored for machine learning experts and specific data types.

There is a lack of similar user-friendly machine learning kits for biologists and bioinformaticians.

2 Solution




We introduce *genomeNLP*

We solve this problem by providing a package which is designed for biological sequence data processing.

Our command line tool requires only the input sequence files and user-defined parameters.

3 Features



Highly visual and open source

Interactive visualisations with plots & tables of metrics and compute resources are generated.

Files are compatible with commonly used tools in the event where low-level customisation is needed.

4 Future



Extend to other methods

We will extend this package continuously with the latest state of the art methods.

Software is open-source and external contributions are welcome at <https://github.com/tyronechen/genomenlp>

Thank you!

Next:

- Hands-on exercise
- Login to Google colaboratory <https://colab.research.google.com>

Classification (\sqrt{p}) vs regression ($p/3$)

1. Compute the accuracy on the i th training set (80-20 split)
2. Compute the accuracy on the j th feature (permuted)
3. Subtract the acc of permuted training set from that of the unmodified training set. The difference will be higher for more important features.
4. Average over all training sets

Text representations are mostly used in,

1. Machine Translation — Automatically translating text from one language to another.
2. Document clustering — Grouping text documents based on the structural and/or semantic similarity.
3. Topic detection — Identifying the topic of a large text corpus.
4. Text summarization.
5. Question Answering.

6.

NLP preprocessing of sequence data:

1. Tokenization: Finding K-mers
2. Numerical encoding:
 - a. Ordinal Encoding
 - b. One-Hot Encoding
3. Transformations
 - a. Frequency tables
 - i. N-grams
 - ii. Term Frequency - Inverse Document Frequency
 - b. Embedding Vector
 - c. Positional encoding

Some of the text representations are [one-hot encoding](#), [n-gram model](#), [Bag-of-Words model](#) and [neural word embedding](#).

In the bag of words model text or the documents are represented by modeling a *bag* (unordered collection) of words. This *bag* of words does not count the positioning, grammar or structure of the words in the text. It just count the frequencies of words in the target text and put that words in to a *bag*. The frequencies of words appearing in a text (sentence or document) is the feature that is used on bag-of-words model.